

Retrieval-based Voice Conversion

The architecture and training of a VITS-derived, GAN-trained retimbre model

A conceptual walkthrough of RVC v2 — grounded in the `rvc` pure-Rust reimplementation

Abstract. — RVC (Retrieval-based Voice Conversion) takes a **source** recording and re-renders it in a **target** voice, keeping **what was said and how it was pitched** while replacing **who is saying it**. This document explains the ideas that make that possible: how the model separates linguistic content and pitch from speaker timbre, how it is structured as a conditional variational autoencoder wrapped around a neural vocoder, and how a **generative adversarial network** is used to teach that vocoder to synthesise waveforms indistinguishable from real audio. Throughout, the concepts are anchored to a concrete open implementation so the abstract picture stays checkable against running code.

Contents

1	Provenance: which RVC this is	2
1.1	“v2” is not a version to be behind on	2
1.2	The warm-start bases	2
1.3	Why 48 kHz is the only rate	3
2	What voice conversion actually is	3
3	Background: a few building blocks	4
4	Analysis: turning audio into content and pitch	5
4.1	ContentVec — a speaker-blind content representation	5
4.2	RMVPE — a robust pitch tracker	5
5	The generative core: a conditional VAE around a vocoder	6
5.1	The two encoders and the latent space	6
5.2	The normalizing flow — bridging prior and posterior	7
6	The vocoder: NSF-HiFiGAN	7
6.1	Why a source-filter model	7
6.2	The deterministic sine source	8
6.3	A note on dimensions (v2, 48 kHz)	8
7	Training I: the reconstruction objectives	9
8	Training II: the generative adversarial network	9
8.1	A tale of two networks	9
8.2	Why many discriminators — multi-period + multi-scale	10
8.3	The four adversarial/perceptual terms, precisely	10
8.4	The alternating step	11
8.5	Reading the loss curves	11
9	Training III: making training stable and reliable	12
9.1	A decaying learning-rate schedule	12
9.2	Averaging the weights: the EMA snapshot	12
9.3	Bigger effective batches for free: gradient accumulation	12

9.4	Keeping the discriminator in check	12
9.5	Weighting the data by cleanliness	12
10	Where training goes wrong: data, not just model	13
11	Inference recap: the reverse path	13

1 Provenance: which RVC this is

Everything below describes **RVC v2**, mirrored from the reference implementation `RVC-Project/Retrieval-based-Voice-Conversion-WebUI` at tag `2.2.231006` — and within it `infer/lib/infer_pack/{models,attentions,modules}.py`, where `SynthesizerTrnMs768NSFsid` and `MultiPeriodDiscriminatorV2` are defined. The Rust modules keep that reference’s **field names** and nesting, so a published PyTorch `state_dict` maps onto the module tree parameter for parameter. That is a standing constraint on the port rather than a convenience: without it there is no warm-start (§7), and on a one-hour corpus there is no useful training without warm-start.

1.1 “v2” is not a version to be behind on

The name invites the assumption that a v3 exists and this port has not caught up. **It does not.** A v3 has been announced in upstream’s release notes since `2.1.230814` (August 2023) and has never shipped. Upstream’s most recent release, `2.3.260718` of 21 July 2026, says so in as many words — **base model unchanged** — and what it does change sits outside the network entirely:

- **packaging and WebUI ergonomics** — installation, defaults, the browser front end. None of it reaches the model.
- **vocal separation swapped from UVR5 to PyMSS** — a corpus-preparation tool that runs before training ever starts. This toolkit uses neither: it ports MDX23C, the TFC-TDF-UNet UVR itself ships, into `burn-mdx`, and its corpus preparation is that plus the sentence-safe slicer of §10.
- **FCPE added as a pitch extractor** — an alternative to RMVPE (§4) in the analysis front end, which is **frozen and speaker-independent** and therefore swappable without touching a single trained weight.
- **CUDA-graph speed-ups on the real-time path** — an inference-scheduling change.

So the tag pinned above fixes the **source text** being mirrored, not the generation of the model: v2 is upstream’s current architecture, and a v2 port is level with it. If v3 ever lands it will be a new network to port, not a migration.

1.2 The warm-start bases

A “base” is a generator/discriminator pair trained on many speakers, which fine-tuning then specialises to one voice (§7). Several exist for **this same architecture**, so they differ only in the weights, never in the shapes of §6:

base	rates	notes
<code>f0G48k.pth</code> / <code>f0D48k.pth</code>	32/40/48 kHz	the stock v2 base, from HF <code>lj1995/VoiceConversionWebUI</code> ; the 48 kHz pair is what this toolkit fetches (76 MB + 143 MB)
TITAN-Medium	32/40/48 kHz	community, Apache-2.0, ships G and D at every rate; trained on 11.15 h of the Espresso corpus
Ov2Super	32/40 kHz	community; no 48 kHz weights published
RIN_E3	32/40 kHz	community; no 48 kHz weights published

Because the architecture is identical, swapping a base is a matter of pointing the trainer’s pretrained generator and discriminator flags at different files — the loader remaps and applies them exactly as it does the stock pair.

TITAN-Medium is the only alternative reachable at 48 kHz, and it is deliberately not wired up. Nothing in the toolkit fetches or selects it, and the reason is that it has not been compared against the stock base on this project’s own corpora; making it a default would ship an untested claim about output quality on soft, breathy material, which is the material that matters here. Passing its G/D pair by hand is expected to work and **has not been tried** — an honest untested path, not a supported one. Ov2Super and RIN_E3 are out of reach for a blunter reason: they publish nothing at 48 kHz, and 48 kHz is the only rate this trainer accepts.

1.3 Why 48 kHz is the only rate

The trainer rejects any other `--model-sr` outright (`crates/rvc-train/src/lib.rs:140-143`). That limit is **not** architectural: `SynthesizerConfig::v2_40k()` already exists (`crates/burn-rvc/src/config.rs:67-76`) and departs from the 48 kHz configuration in two lines — upsample rates $10 \cdot 10 \cdot 2 \cdot 2$ instead of $12 \cdot 10 \cdot 2 \cdot 2$, giving a hop of 400 samples rather than 480, with transposed-conv kernels to match. Everything **around** the network is what assumes 480: the analysis frame grid, the training STFT front end ($n_{\text{fft}} = 2048$, hop 480), the dataset windowing, and the base weights that would have to be fetched at the new rate. Reaching 40 kHz is therefore a plumbing job rather than a modelling one — and until someone does it, the two 32/40 kHz community bases cannot be used at all.

2 What voice conversion actually is

A speech signal braids together several independent things at once: the **linguistic content** (the sequence of phonemes — the words), the **prosody** (the pitch contour, or fundamental frequency F_0 , and rhythm), and the **timbre** (the resonances of a particular vocal tract — the thing that makes a voice recognisable). Text-to-speech generates all three from scratch. Voice conversion is a narrower, in some ways harder, problem: keep content and prosody **exactly**, and swap only the timbre.

The design principle that follows is **disentanglement**. If we can extract a representation of a frame of audio that captures **what is being said** but is **blind to who is saying it**, then we can hand that content-only representation, together with the desired pitch, to a decoder that has learned one specific target voice — and the decoder has no choice but to render the same content in that voice. RVC leans on this hard, and it is why breathy, non-lexical vocalizations (sighs, gasps) survive conversion by construction: they are part of the content and pitch streams, not something the model has to “recognise.”




Figure 1: The end-to-end conversion path. Everything above the decoder is speaker-independent analysis; the decoder is the only speaker-specific part.

The rest of this document unfolds that diagram top to bottom, then explains how the speaker-specific decoder is **trained** — which is where the GAN enters.

3 Background: a few building blocks

Readers new to audio machine learning may want these six ideas first; every later section leans on them. Experienced readers can skip ahead.

Digital audio — samples & sample rate a sound wave is stored as a long list of numbers, the **samples**, measured many thousands of times per second. The **sample rate** is how many per second. Analysis here runs at 16 kHz (16 000 samples/s — enough to read content and pitch), while the target voice is synthesised at 48 kHz for full fidelity.

Spectrogram (STFT) the raw sample list is hard to reason about directly. Sliding a short window along the signal and taking a **Fourier transform** of each window gives the **short-time Fourier**

transform — a 2-D picture of **how much energy sits at each frequency, moment by moment**. Columns are time frames; rows are frequency bins.

The mel scale & mel spectrogram human hearing resolves low frequencies far more finely than high ones. The **mel scale** warps frequency to match, $\text{mel}(f) = 1127 \cdot \ln\left(1 + \frac{f}{700}\right)$. A **mel spectrogram** is an STFT whose frequency rows have been regrouped into a smaller set of perceptually spaced **mel bands** (128 of them here). Because closeness on a mel spectrogram tracks “sounds alike,” the main training loss is simply an L_1 distance between the mel spectrograms of the real and generated audio (§7).

Embeddings a network cannot consume a raw category (“speaker 3”, “pitch bin 57”) directly. An **embedding** is a learned lookup table turning each discrete id into a short vector the network can use; because it is learned, similar ids drift to similar vectors. RVC embeds both the **speaker** (g) and the **coarse pitch**.

Convolutions a **convolution** slides a small learnable filter along the signal and fires wherever its local pattern appears. Stacking them detects increasingly complex structure (edges $\rightarrow$ harmonics $\rightarrow$ textures) while reusing the same weights everywhere — efficient and natural for audio, where the same waveform shapes recur throughout. Both the vocoder and the discriminators are convolutional.

Training — loss & gradient descent a **loss** is one number measuring how wrong the output is.

Gradient descent nudges every weight a tiny step in the direction that most reduces it; repeated over many examples, the network learns. The step size is the **learning rate**, and how it is scheduled turns out to matter for stability (§9).

4 Analysis: turning audio into content and pitch

Two pretrained, frozen models do the analysis. Neither is trained by RVC; they are universal front-ends, run identically at training and inference time (in this implementation they are shared ONNX sessions behind one `FeatureExtractor`).

4.1 ContentVec — a speaker-blind content representation

ContentVec is a **self-supervised** speech encoder (a HuBERT-family transformer) that maps 16 kHz audio to a sequence of 768-dimensional vectors, roughly one per 20 ms frame. Its crucial property for our purposes is that it was trained with a speaker-disentanglement objective: two people saying the same words produce nearly the same ContentVec sequence. That vector is therefore a good proxy for “linguistic content, timbre removed” — exactly the input a voice converter wants.

In code the content stream is extracted at 50 Hz and **upsampled $\times 2$** to the 100 Hz frame grid so it aligns one-to-one with the pitch stream:

```
let up = upsample_rows(&feats.content, 2); // 50 Hz -> 100 Hz
let n = up.len().min(feats.f0.len()); // align to F0 length
```

4.2 RMVPE — a robust pitch tracker

RMVPE estimates the fundamental frequency F_0 — the pitch — for each frame, and marks frames as **voiced** or **unvoiced** (silence / breath has no pitch). It is prized for staying accurate on noisy, expressive, or quiet material, which matters enormously for quiet, close-mic corpora full of soft phonation.

Pitch feeds the model in **two** forms, and understanding why is worth a moment:

Coarse pitch (integer ids [1,255]) the continuous F_0 in Hz is warped onto a **mel scale** and quantized into 255 bins. This becomes an **embedding lookup** inside the prior encoder — it is treated like a discrete symbol, letting the network learn a smooth notion of “pitch class.”

Fine F_0 (Hz, continuous) the raw value drives the decoder’s harmonic oscillator directly, so the synthesised waveform sits at **precisely** the right pitch, not the nearest bin.

The mel-warping quantization (from `dsp::f0_to_coarse`) is:

$$\text{mel}(f) = 1127 \cdot \ln\left(1 + \frac{f}{700}\right), \quad \text{id} = \text{round}(\text{clamp}(1..255)).$$

Unvoiced frames ($f = 0$) map to id 1; the map is monotonic, so higher pitch always yields a higher (or equal) bin.

5 The generative core: a conditional VAE around a vocoder

RVC’s network is a descendant of **VITS**, adapted for conversion and named **SynthesizerTrnMs768NSFsid**. It is best understood as a **conditional variational autoencoder** (CVAE) whose decoder happens to be a full neural vocoder. Five submodules cooperate:




Figure 2: The five modules of the synthesizer and the two data paths. The posterior encoder (dashed) exists **only during training** — it is the “cheat sheet” that teaches the prior what a good latent looks like.

5.1 The two encoders and the latent space

Both encoders emit a **distribution**, not a point: a per-frame mean μ and log-standard-deviation $\log \sigma$ over a 192-dimensional latent. This is the VAE idea — the latent z is stochastic, sampled with the reparameterization trick $z = \mu + \varepsilon \cdot e^{\log \sigma}$, $\varepsilon \sim \mathcal{N}(0, 1)$.

- The **posterior encoder enc_q** sees the **real answer**: the linear spectrogram of the ground-truth target waveform. It is a stack of 16 dilated WaveNet layers and produces the posterior $(\mu_q, \log \sigma_q)$. Because it gets the actual audio, its latent z is easy for the decoder to turn back into that audio. But at inference we do **not** have the target audio — that is the whole point — so **enc_q** cannot be used then.
- The **prior encoder enc_p** sees only **content + pitch** — the information we **will** have at inference. It is a 6-layer relative-position transformer producing the prior $(\mu_p, \log \sigma_p)$.

The trick that ties them together is a **normalizing flow**.

5.2 The normalizing flow — bridging prior and posterior

There is a gap: the posterior latent z (rich, from real audio) and the prior z_p (from content only) live in different-looking distributions. The flow is an **invertible** neural function f_g (conditioned on the speaker g) that warps between them, and invertibility is what lets us run it **both directions**:




Figure 3: The flow is one function used two ways. Training pushes the posterior latent toward the prior; inference pulls a prior sample into the decoder’s latent space.

RVC’s flow is a **ResidualCouplingBlock** of four **additive coupling layers**. Each layer splits the channels in half, leaves the first half untouched, and adds to the second half a WaveNet-predicted shift computed **from the first half**:

$$x_1 \leftarrow x_1 + m(x_0, g) \quad (\text{forward}), \quad x_1 \leftarrow x_1 - m(x_0, g) \quad (\text{reverse}).$$

Because the shift depends only on the **untouched** half, the layer is trivially invertible — reverse just subtracts the very same predicted mean. A channel **flip** between layers ensures every dimension eventually gets transformed. The layers are **mean_only**, so the log-determinant is zero and there is nothing extra to track:

```

pub fn forward(&self, x, g) { cat(x0, x1 + mean(x0, g)) } // add
pub fn reverse(&self, x, g) { cat(x0, x1 - mean(x0, g)) } // subtract
  
```

At inference the sampling temperature is deliberately cooled — $z_p = \mu_p + \sigma_p \cdot \varepsilon \cdot 0.666$ — trading a little diversity for cleaner, more stable output.

6 The vocoder: NSF-HiFiGAN

The decoder **dec** turns the 192-channel latent (at 100 Hz) into a waveform at 48 kHz. That is a $480\times$ increase in sample rate — the **hop length**, the product of the upsample factors $12 \cdot 10 \cdot 2 \cdot 2 = 480$. It is a **HiFi-GAN** generator with a **Neural Source-Filter** (NSF) twist.

6.1 Why a source-filter model

Classic HiFi-GAN upsamples the latent through transposed convolutions and lets the network invent all the fine structure. For **singing and expressive speech** that struggles: the network has to hallucinate an exact pitch. NSF fixes this by **handing the network the pitch as an explicit excitation signal**. The **SourceModule** builds a sine wave whose instantaneous frequency tracks the fine F_0 contour — a clean harmonic **source** — and injects it into every upsampling stage. The convolutional stack then acts as a learned **filter** shaping that source into a natural voice. Pitch accuracy comes for free because it is built into the source.




Figure 4: Inside the NSF-HiFiGAN decoder. The sine source (bottom) is generated deterministically from F_0 and mixed in at each upsampling stage; the speaker embedding g conditions the head.

6.2 The deterministic sine source

The sine generator is **not learned** (RVC runs it under `no_grad`; here it is plain arithmetic). For each frame it computes a phase increment f_0/sr , accumulates phase across the whole clip, upsamples the cumulative phase (linear, `align_corners`) and takes its sine — carefully detecting phase wraps so the oscillator stays continuous:

```
phase += rad_up[j] + shift;           // shift = -1 on a detected wrap
sine[j] = (phase * 2.0 * PI).sin() * SINE_AMP;
```

Voiced frames get the sine; unvoiced frames get low-level Gaussian noise (breath). The only learnable part of the source is a single `l_linear` projection feeding a `tanh`. This is precisely why **breathy, unvoiced speech survives**: unvoiced frames are handled explicitly by the noise branch, not discarded.

6.3 A note on dimensions (v2, 48 kHz)

quantity	value	meaning
content dim	768	ContentVec output width, per frame
latent <code>inter_channels</code>	192	shared width of prior / posterior / flow
<code>hidden_channels</code>	192	transformer & WaveNet width
<code>spec_channels</code>	1025	linear-spectrogram bins ($\frac{n_m}{2} + 1$) into <code>enc_q</code>
<code>gin_channels</code>	256	speaker-embedding width g
transformer	6 layers, 2 heads	prior encoder <code>enc_p</code>
flow	4 coupling layers	additive, <code>mean_only</code> , with flips
upsample rates	$12 \cdot 10 \cdot 2 \cdot 2$	product = hop = 480 samples/frame
frame rate	100 Hz	both content and pitch land on this grid

7 Training I: the reconstruction objectives

Training is **fine-tuning**: warm-start every module from the public pretrained base (`f0G48k.pth` / `f0D48k.pth`), then adapt to one target voice on a small corpus. The module layout is deliberately kept weight-compatible with the reference PyTorch `state_dict` so this warm-start is possible — essential when the corpus is only 1 hour long. Other bases exist for the same architecture; §1 lists them and says why none of them is wired up.

Each step draws a batch of short random windows from the corpus, runs the **training forward** (posterior $\rightarrow$ flow $\rightarrow$ decode a random 0.36 s segment), and computes a compound loss. Three of its terms are **reconstruction** terms; the other two (next section) are **adversarial**.

Mel-spectrogram L1 ($\times 45$) the perceptual backbone. Convert both the real segment and the generated one to a 128-band mel spectrogram and take the mean absolute difference. This is what actually makes the output **sound** like the target; its large weight (45) reflects that.

$$L_{\text{mel}} = \| \text{mel}(y) - \text{mel}(\hat{y}) \|_1 .$$

KL divergence ($\times 1$) the CVAE glue. It pushes the flow-transformed posterior z_p to look like a sample from the prior $(\mu_p, \log \sigma_p)$, so that at inference — when only the prior is available — the decoder still receives a latent it knows how to render. RVC uses VITS’s Monte-Carlo form:

$$L_{\text{kl}} = \log \sigma_p - \log \sigma_q - \frac{1}{2} + \frac{1}{2} (z_p - \mu_p)^2 e^{-2 \log \sigma_p} .$$

Feature matching ($\times 2$) defined via the discriminator, covered below.

A subtle but important detail: the posterior’s input spectrogram is `.detach()`-ed, and in feature matching the **real** features are detached — they are fixed targets, and gradients must not flow into them.

8 Training II: the generative adversarial network

Mel-L1 alone produces **oversmoothed**, **buzzy** audio: averaging many plausible waveforms that share a mel spectrogram gives a muddy one. The fix is a **GAN**. We introduce a second network — the **discriminator** — whose only job is to tell **real** target audio from the **generated** audio. The generator (our whole synthesizer) is then additionally trained to **fool** it. In the limit, “audio the discriminator can’t distinguish from real” is exactly “audio that has the crisp fine structure of real speech.”

8.1 A tale of two networks




Figure 5: The adversarial game. G wants D to score its output as real; D wants to score real as 1 and fake as 0. They are trained in alternation, each with its own AdamW optimiser.

8.2 Why many discriminators — multi-period + multi-scale

A single discriminator looking at the raw 1-D waveform tends to miss **periodic** artefacts — the buzz and roughness that live in the pitch harmonics. RVC (from HiFi-GAN) therefore uses a **bank** of discriminators, a `MultiPeriodDiscriminatorV2`:

- One **scale** discriminator (`DiscriminatorS`) reads the waveform as a plain 1-D signal — good at broad envelope and noise texture.
- Eight **period** discriminators (`DiscriminatorP`), one for each prime period in $\{2, 3, 5, 7, 11, 17, 23, 37\}$. Each **folds** the 1-D waveform into a 2-D grid with that period as the row stride, then runs 2-D convolutions. Folding on a period exposes structure that repeats at that spacing — precisely where pitch-related artefacts hide. Using **coprime** periods means the discriminators cover a wide range of periodicities without redundantly overlapping.




Figure 6: How a period discriminator “sees” the signal: a period-3 fold turns a 1-D strip into a 2-D image whose columns line up periodic samples, then convolves it in 2-D. Eight coprime periods cover many spacings.

Both discriminator families also expose their **intermediate feature maps**, which powers the feature-matching loss.

8.3 The four adversarial/perceptual terms, precisely

RVC uses the **least-squares** GAN (LSGAN) formulation — it is more stable than the log-loss original. Writing D_k for the k -th discriminator’s score:

Discriminator loss (push real→1, fake→0)

$$L(D) = \sum_k \mathbb{E}[(D_k(y) - 1)^2] + \mathbb{E}[D_k(\hat{y})^2].$$

Generator adversarial loss (push fake→1)

$$L_{\text{adv}(G)} = \sum_k \mathbb{E}[(D_k(\hat{y}) - 1)^2].$$

Feature matching ($\times 2$) match G ’s audio to the target **inside** the discriminator, layer by layer — a perceptual L_1 that stabilises training:

$$L_{\text{fm}} = \sum_k \sum_l \left\| D_k^{(l)}(y) - D_k^{(l)}(\hat{y}) \right\|_1.$$

The complete generator objective is the weighted sum

$$L_G = L_{\text{adv}(G)} + 2L_{\text{fm}} + 45L_{\text{mel}} + L_{\text{kl}}.$$

8.4 The alternating step

One training step touches both networks. For the **discriminator** term the generated audio is **detached**, so D 's gradient never leaks into G . Both the D and the G gradients are computed from the **same start-of-step weights**, and only **then** are the two optimisers stepped — so within a step G is pushed against the discriminator as it stood at the start, not a half-updated one. (With gradient accumulation, §9, these gradients are summed over several micro-batches before the step.)




Figure 7: One training iteration. Both losses are measured against the same start-of-step weights; the two optimisers then step together.

Both optimisers are AdamW ($\beta_1\{=\}0.8$, $\beta_2\{=\}0.99$), matching the reference recipe, with a base learning rate of 10^{-4} that is **decayed over the run** (§9). What ships is not the raw final weights but a smoothed **exponential moving average** of them — also §9.

8.5 Reading the loss curves

Because G and D are **adversaries**, their losses do **not** both fall to zero — they seek an equilibrium. Practical intuition:

- **mel_loss** is the one to watch. It is the honest reconstruction signal and **should trend down**. Some **shaking around a plateau** in the back half is normal — it is the adversarial equilibrium, and the stabilisation techniques of §9 (LR decay, weight EMA) exist to tame exactly that. But if it **never** descends and the output is silent, suspect a **data** problem (§10), not a hyperparameter one.
- **d_loss** near a **small positive constant** (not collapsing to 0) means the discriminator is appropriately challenged. A **d_loss** crashing to 0 means D has won and G gets no useful gradient.
- **g_loss** is dominated by the $45 \times \text{mel}$ term, so it largely tracks **mel_loss** plus adversarial jitter.

9 Training III: making training stable and reliable

The recipe so far — LSGAN, feature matching, mel-L1, KL — is faithful to the reference. But an adversarial game trained on a **small** corpus and a **small** GPU is a jittery thing: with a tiny batch each step’s gradient is noisy, and once G and D settle into their tug-of-war the mel loss stops descending and instead **oscillates** around a plateau. Left alone, whatever noisy point the final step lands on is what gets saved. Five techniques — all optional `rvc train` flags, two on by default — make the outcome steadier and the saved model better. None of them touches the **objective**; they shape the **dynamics**.

9.1 A decaying learning-rate schedule

A constant learning rate keeps taking full-size steps forever, so near the optimum the weights **bounce around** it instead of settling in. The remedy is to shrink the step as training proceeds. Here the rate decays **exponentially over the whole run**,

$$\text{lr}(t) = \text{lr}_0 \cdot \rho^{t/T},$$

from lr_0 (the `--lr` base) at the first step down to $\text{lr}_0 \cdot \rho$ at the last, where $\rho = \text{--lr-final}$ and T is the total number of steps. Tying the schedule to the **fraction of the run completed** rather than to an epoch count makes it behave identically whether you train for two epochs or forty.

9.2 Averaging the weights: the EMA snapshot

Even with a decaying rate the weights still wobble step to step. A cheap, standard trick from GAN vocoders is to keep a second, **shadow** copy of the generator that trails the live weights — an **exponential moving average**:

$$\theta_{\text{ema}} \leftarrow d \cdot \theta_{\text{ema}} + (1 - d) \cdot \theta_{\text{live}} \quad (\text{every step}).$$

Averaging over the recent past cancels the oscillation, so the EMA weights are **cleaner and less buzzy** than any single step. Two points worth stressing: the EMA is a **read-only snapshot** — it never feeds back into the optimiser, so it does **not** slow learning — and it is what gets **saved as the model** (the raw live weights are written beside it, so nothing is lost, and on a very short run that never plateaus you can prefer them). The averaging window is set as a fraction of the run (`--ema-frac`), so, like the LR schedule, it needs no re-tuning when the run length changes.

9.3 Bigger effective batches for free: gradient accumulation

A 6 GB GPU forces a tiny batch, and tiny batches make noisy gradients. Gradient **accumulation** recovers a larger **effective** batch without more memory: run `--grad-accum N` micro-batches, **sum** their gradients, and step once. Peak memory stays that of a single micro-batch — each one’s computation graph is freed after its backward pass — while the summed gradient is as steady as a batch N times larger, at $N \times$ the compute.

9.4 Keeping the discriminator in check

If D learns much faster than G it starts winning outright: its gradients stop being informative and instead inject **buzzy, high-frequency artefacts** into the generator. Two knobs rebalance the game — `--d-lr-ratio` scales D ’s learning rate below G ’s, and `--d-interval` updates D only every few steps — both handing G a little room to catch up.

9.5 Weighting the data by cleanliness

Finally, corpus clips are not equally clean. `--snr-weight` biases sampling toward the clips with a lower noise floor, weighting each by SNR^α . Crucially the score is a **signal-to-noise ratio** — a clip’s loud-percentile level over its noise-floor level — and **not loudness**, so a soft, breathy take still scores

high and is never penalised for being quiet. This steers away from hiss while keeping the very content the corpus exists to capture. (Contrast §10, which removes between-sentence **silence**; this weights whole clips by **noise**.)

10 Where training goes wrong: data, not just model

The most common failure on real-world corpora is a generator that **collapses to silence** — `mel_loss` stuck and shaking, output pure silence, even though the pretrained base works. The architecture is fine; the **training distribution** is poisoned.

The trainer draws **random short windows uniformly across each file**. If the raw recordings are full of between-sentence dead air, then most randomly sampled 0.36 s segments **are silence**, and the fastest way to minimise mel-L1 on a silence-heavy batch is to **emit silence**. The generator obliges, globally.

The remedy is a **sentence-safe slicer** run as a preprocessing pass (`preprocess clip` , in the binary of that name — it was `rvc preprocess` until corpus preparation became an engine-shaped pair of its own): it uses short-time energy with **hysteresis** to cut only **genuine, sustained** dead air between sentences, while:

- **never splitting a complete sentence** — a gap must be both below the energy floor **and** longer than a minimum duration before it counts as a cut; and
- **preserving soft, breathy content** — energy is used only to **locate** long silent gaps, never to gate quiet-but-present sound, and cut boundaries are edge-padded so onsets and breathy tails are kept.

The result is a corpus of clean per-sentence clips, so uniform window sampling now lands on real phonation. Training itself is unchanged — it simply consumes a healthier distribution.

11 Inference recap: the reverse path

At inference the posterior encoder and the whole discriminator bank are gone. The path collapses to: analyse source → prior encoder → **sample** z_p → **reverse** the flow → decode with the target speaker embedding. Content and pitch came from the **source**; timbre came from the **trained decoder + speaker embedding**. That asymmetry — universal analysis in, speaker-specific synthesis out — is the whole idea of voice conversion, and every architectural choice above exists to keep those two halves cleanly separated.

Concepts anchored to the `rvc` implementation: `burn-rvc` (network), `rvc-core` (feature extraction + DSP + backends), `rvc-train` (adversarial training loop). RVC v2, 48 kHz, weight-compatible with the reference `SynthesizerTrnMs768NSFsid` / `MultiPeriodDiscriminatorV2` (RVC-Project, tag `2.2.231006`).